


Given examples, astra solves most tasks ...

gpt-6-astra on P-ARC (50 tasks) · official test item correct, % · mean of two runs



... but as an active discoverer, choosing its own examples, it falls far behind