


Given examples, state-of-the-art models solve most tasks ...

Official test item correct, % · mean of two runs



... but as active discoverers, choosing their own examples, they fall far behind